



Figure 1: A high-level view of autotuning methods, where **blue** boxes denote branching questions, **orange** boxes denote key hypotheses, **green** boxes denote algorithm choices, and **bold** boxes denote methods.